



Flash - Cloudera Impala on Docker

Big Data Montreal #29

October 7th, 2014



Tony Bussieres, SVP Dev & co-founder

tony.bussieres@ticksmith.com

Running Impala in a Docker container



<http://impala.io>



<http://docker.com>

What is Docker

- A portable and lightweight application container.
- It's not a VM
- Open source, works on most modern OS.
- Images are easy to create and maintain with Dockerfiles
- Images are easy to share with developers / sysadmin using the Docker Hub
- More info on Docker at <https://docs.docker.com/>
- **It's not a VM!** (It shares the same kernel & resources as the host)

Motivation

- It makes it easy to use Impala on your desktop
- Test queries without a connection to a cluster
- More lightweight than using Cloudera's QuickStart VM
- Wanted to prove that Impala could run in Docker
- Docker is cool ! :)

Quick Demo

Docker image **codingtony/impala** is based on CDH5.
It is configured as a Single-Node Hadoop cluster

It runs a namenode, a datanode, the impala-state-store, the impala-catalog and the impala-server

The Docker Image is available on Docker Hub :
docker pull codingtony/impala

This presentation is available :
<https://github.com/codingtony/docker-impala-bdm29>

The recipe for building **codingtony/impala** is available here :
<https://github.com/codingtony/docker-impala>



Any Questions ???